



# **Optimizing Blackjack Strategies**

Clarissa Nunez, Jacob Fischbach,  
Jacob Klucher, James Elander &  
Luke Milton



# Project Overview

## Goal

Develop and analyze optimized strategies for playing Blackjack using machine learning algorithms.

## Specifics

Use various techniques to build a model that predicts the best possible move based on the player's hand, the dealer's upcard, and the game state.

## Limitations

Due to potential issues from complexity, the model will not split, double, insurance, or use card counting.

# Technologies Used

Machine  
Learning

**Keras & TensorFlow**

Used to create the  
Actor Critic Default  
Program



Python

**Python**

Used to create the  
Blackjack game



Matplotlib

**Matplotlib**

Used to create plots  
for model  
predictions




# The Process

Create a Blackjack environment

Use the Blackjack environment to train the Actor-Critic algorithm

The Actor-Critic algorithm should be able to predict the best action for the player: hit or stay



Construct an Actor-Critic Reinforcement Learning Network

Optimize the decision-making process for Blackjack players


# Blackjack Environment

```
1 import random
2
3 # Card values (Ace can be 1 or 11)
4 CARD_VALUES = {
5     "2": 2, "3": 3, "4": 4, "5": 5, "6": 6, "7": 7, "8": 8, "9": 9, "10": 10,
6     "J": 10, "Q": 10, "K": 10, "A": 11
7 }
8
9 # Function to calculate hand value
10 def calculate_hand_value(hand):
11     value = sum(CARD_VALUES[card] for card in hand)
12     aces = (variable) value: int
13     while value > 21 and aces:
14         value -= 10 # Convert Ace from 11 to 1
15         aces -= 1
16     return value
17
18 # Function to deal a card
19 def deal_card():
20     return random.choice(list(CARD_VALUES.keys()))
21
22 # Function to play a round of blackjack
23 def play_blackjack():
24     player_hand = [deal_card(), deal_card()]
25     dealer_hand = [deal_card(), deal_card()]
26
27     print(f"Your hand: {player_hand}, Value: {calculate_hand_value(player_hand)}")
28     print(f"Dealer's first card: {dealer_hand[0]}")
29
30     # Player's turn
31     while calculate_hand_value(player_hand) < 21:
32         action = input("Do you want to hit or stay? (h/s): ").lower()
33         if action == 'h':
34             player_hand.append(deal_card())
35             print(f"Your hand: {player_hand}, Value: {calculate_hand_value(player_hand)}")
36         else:
37             break
```

```
39     player_value = calculate_hand_value(player_hand)
40     if player_value > 21:
41         print("You busted! Dealer wins.")
42         return
43
44     # Dealer's turn
45     print(f"Dealer's hand: {dealer_hand}, Value: {calculate_hand_value(dealer_hand)}")
46     while calculate_hand_value(dealer_hand) < 17:
47         dealer_hand.append(deal_card())
48         print(f"Dealer hits: {dealer_hand}, Value: {calculate_hand_value(dealer_hand)}")
49
50     dealer_value = calculate_hand_value(dealer_hand)
51     if dealer_value > 21:
52         print("Dealer busted! You win!")
53     elif dealer_value > player_value:
54         print("Dealer wins.")
55     elif dealer_value < player_value:
56         print("You win!")
57     else:
58         print("It's a tie!")
59
60 # Run the game
61 if __name__ == "__main__":
62     while True:
63         play_blackjack()
64         again = input("Play again? (y/n): ").lower()
65         if again != 'y':
66             break
```

# Actor-Critic Explained

- Actor-Critic reinforcement learning combines policy-based and value-based methods
- ACs use two neural networks: an “actor” that selects actions and a “critic” that evaluates those actions, by estimating their value or quality
- Actor: makes decisions by selecting actions based on current policy, it's goal is to maximize rewards by exploring the action space
- Critic: evaluates the actions taken by the actor, it estimates the value/quality and provides feedback on the actor's performance







# Actor-Critic Algorithm

## Two Inputs

- Player's hand value
- Dealer's visible card

## Two Hidden Layers

- Two dense layers with ReLU activation for processing input data

## Action Prediction

- Actor - The model outputs probabilities of the actions (Hit or Stay) and chooses the action based on either the probabilities or by exploring

## Value Estimation

- Critic - The model predicts the expected future reward for the given state

The background of the slide features a blurred image of playing cards, specifically the Ace of Diamonds and the Ace of Hearts, along with a small cartoon character of a person in a red hat and blue outfit.

# Actor-Critic Algorithm

At each step, the model observes the current state of the game and selects an action (Hit or Stay) based on the actor's output.

The model receives a reward after each action

- Win: +1
- Loss: -1
- Tie: 0

Then it uses these rewards to update its parameters

- Actor Loss - How well the actor chose the action
- Critic Loss - How accurate the critic's value estimation was



The background of the slide features a blurred image of playing cards, specifically the Ace of Diamonds and the Ace of Hearts, along with a small, stylized character on the left side.

# Actor-Critic Algorithm

```
# Setup
import keras
import tensorflow as tf
from keras import layers
import numpy as np
from blackjack_env import BlackjackEnv
import matplotlib.pyplot as plt
import pandas as pd
import random

# Custom blackjack environment
env = BlackjackEnv()

# Configuration parameters for the whole setup
gamma = 0.99 # Discount factor for past rewards
num_inputs = 2 # Player hand value, Dealer's visible card
num_actions = 2 # Hit or stay
num_hidden1 = 64 # 8, 16, 32, 64, 128, 256
num_hidden2 = 64

# Set random seed
seed_value = 42
np.random.seed(seed_value)
random.seed(seed_value)
tf.random.set_seed(seed_value)
```

# Actor-Critic Algorithm

```
# Define the Actor-Critic Model
# Input
inputs = layers.Input(shape=(num_inputs,))

# Hidden
hidden1 = layers.Dense(num_hidden1, activation="relu")(inputs)
hidden2 = layers.Dense(num_hidden2, activation="relu")(hidden1)

# Actor
action = layers.Dense(num_actions, activation="softmax")(hidden2)

# Critic
critic = layers.Dense(1)(hidden2)

# Create model
model = keras.Model(inputs=inputs, outputs=[action, critic])
#####
initial_learning_rate = 0.005
lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(
    initial_learning_rate, decay_steps=100, decay_rate=0.96, staircase=True
)
optimizer = tf.keras.optimizers.Adam(learning_rate=lr_schedule)
#####
#optimizer = keras.optimizers.Adam(learning_rate=0.005) # Different rates: 0.005, 0.001 , 0.0005, 0.0001
huber_loss = keras.losses.Huber()
```

# Actor-Critic Algorithm

```
epsilon = 0.1 # Exploration factor
epsilon_min = 0.01
epsilon_decay = 0.995 # 0.999, 0.995, 0.99, 0.95
```

```
for hand_count in range(MAX_HANDS):
    state = env.reset()
    episode_reward = 0
    action_probs_history, critic_value_history, rewards_history = [], [], []

    with tf.GradientTape() as tape:
        for _ in range(100): # Play rounds
            state_tensor = tf.convert_to_tensor(state)
            state_tensor = tf.expand_dims(state_tensor, 0)

            # Predict action probabilities and estimated future rewards
            action_probs, critic_value = model(state_tensor)
            critic_value_history.append(critic_value[0, 0])

            # Epsilon decay
            if epsilon > epsilon_min:
                epsilon *= epsilon_decay

            # Epsilon-greedy action selection
            if np.random.rand() < epsilon:
                action = np.random.choice(num_actions) # Explore
            else:
                action = np.argmax(action_probs) # Exploit

            # Choose an action based on probabilities
            # action = np.random.choice(num_actions, p=np.squeeze(action_probs))
            action_probs_history.append(tf.math.log(action_probs[0, action]))
```

# Actor-Critic Algorithm

```
# Apply the action in Blackjack
state, reward, done, player_value, dealer_value = env.step(action)
#####
if reward == 1:
    won = 1
elif reward == -1:
    won = -1
else:
    won = 0

if player_value >= 17 and reward == 1:
    reward *= 1.1
if player_value >= 22 and reward == -1:
    reward *= 1.1
#####
rewards_history.append(reward)
episode_reward += reward
```

```
Initial Actor Loss: -1.0488
Final Actor Loss: -0.4770
Initial Critic Loss: 1.4647
Final Critic Loss: 0.5615
Average Actor Loss: -0.1528
Average Critic Loss: 0.2477
Average Reward after 2000 episodes: -0.20
Total Wins: 738
Total Losses: 1162
Total Ties: 100
```




# Optimizing the Model:


- Two hidden layers
  - Two layers ended working the best for the model
  - 64 neurons it became more unstable the higher it was
- Learning\_rate = 0.001
- Epsilon-Greedy:
  - epsilon = 0.2
  - epsilon\_min = 0.01
  - epsilon\_decay = 0.995



# Actor-Critic Loss: 100,000

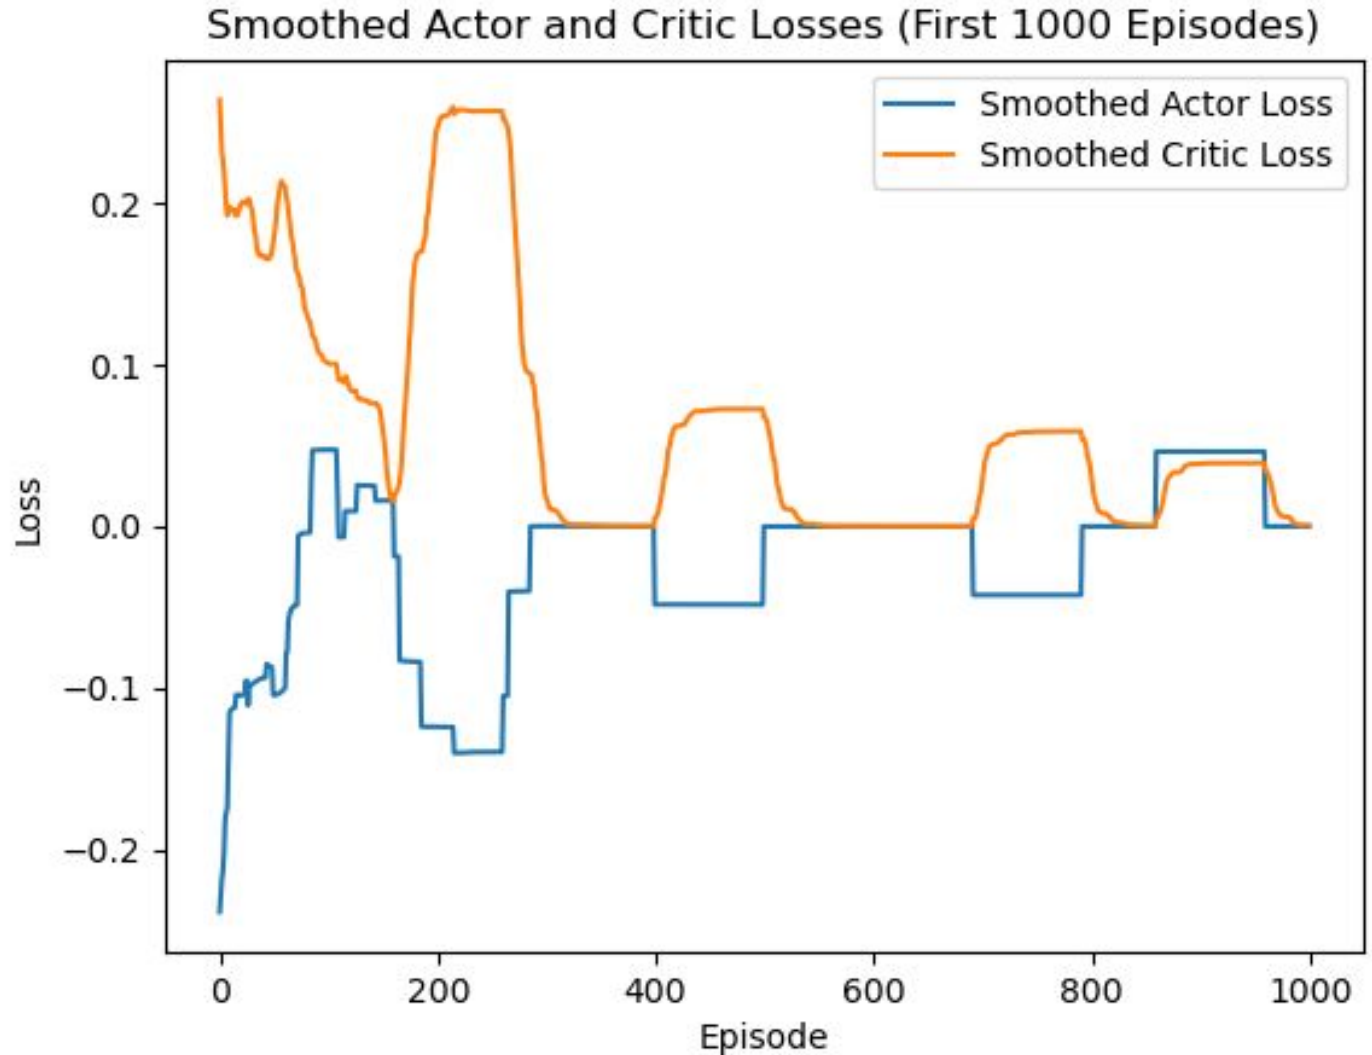


# Actor-Critic Loss: 10,000






# Actor-Critic Loss: 1,000



# Win Rate



# Conclusion

- Created a Blackjack environment in Python
- Created an Actor-Critic Algorithm
- The Actor-Critic Algorithm informs the player whether to hit or stay based on the information available
- Our group was able to optimize the Actor-Critic as best we could within the time constraint
- Next Steps:
  - More actions and more inputs, which likely would increase accuracy
  - Ex: having an Ace in your hand, counting cards, or further hyper tuning the model



**Questions?**